

Final Project on Breast Cancer Wisconsin Dataset:

Breast cancer is one of the most common and life-impacting diseases worldwide, and early detection plays a crucial role in improving patient outcomes. Machine learning has become a powerful tool in medical diagnostics, offering the ability to identify patterns in clinical data that may not be immediately visible to the human eye. In this project, we apply supervised learning techniques to the Breast Cancer Wisconsin Diagnostic Dataset, a widely used benchmark dataset in machine learning and healthcare research.

The objective of this project is to build predictive models capable of distinguishing between malignant and benign breast tumors based on features computed from digitized images of fine-needle aspirates (FNA) of breast masses. These features describe various properties of cell nuclei, such as radius, texture, smoothness, and concavity.

This project follows a full machine-learning workflow:

Problem Definition: - I aim to create a classification model that predicts tumor type (malignant or benign) from the provided numeric features.

Exploratory Data Analysis (EDA): - I analyze the structure of the dataset, examine feature distributions, evaluate correlations, and visualize the data using pairplots and PCA.

Model Building & Training: - Several supervised learning models are trained, including Logistic Regression, Random Forest, Support Vector Machine (SVM) we studied at class.

Model Evaluation & Comparison: - Models are evaluated using accuracy, classification reports, and cross-validation. Hyperparameter tuning via GridSearchCV is used to optimize performance.

Conclusion & Insights: - I identify which model performs best and discuss why it is suitable for this dataset, along with key takeaways about dataset characteristics and model behavior.

This project demonstrates how classical machine learning models can be applied to a real-world medical classification problem, highlighting the importance of data preprocessing, model selection, and careful evaluation. Ultimately, the goal is to build an accurate, interpretable, and computationally efficient model for breast cancer diagnosis.

```
In [1]: from sklearn.datasets import load_breast_cancer
import pandas as pd

# Load dataset
data = load_breast_cancer()

# Create DataFrame
df = pd.DataFrame(data.data, columns=data.feature_names)
df['target'] = data.target

df.head()
```

Out[1]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean di
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	

5 rows × 31 columns

Exploratory Data Analysis (EDA):

In [2]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   mean radius      569 non-null    float64
 1   mean texture     569 non-null    float64
 2   mean perimeter   569 non-null    float64
 3   mean area        569 non-null    float64
 4   mean smoothness  569 non-null    float64
 5   mean compactness 569 non-null    float64
 6   mean concavity   569 non-null    float64
 7   mean concave points 569 non-null    float64
 8   mean symmetry    569 non-null    float64
 9   mean fractal dimension 569 non-null    float64
 10  radius error     569 non-null    float64
 11  texture error    569 non-null    float64
 12  perimeter error  569 non-null    float64
 13  area error       569 non-null    float64
 14  smoothness error 569 non-null    float64
 15  compactness error 569 non-null    float64
 16  concavity error  569 non-null    float64
 17  concave points error 569 non-null    float64
 18  symmetry error   569 non-null    float64
 19  fractal dimension error 569 non-null    float64
 20  worst radius     569 non-null    float64
 21  worst texture    569 non-null    float64
 22  worst perimeter   569 non-null    float64
 23  worst area        569 non-null    float64
 24  worst smoothness  569 non-null    float64
 25  worst compactness 569 non-null    float64
 26  worst concavity   569 non-null    float64
 27  worst concave points 569 non-null    float64
 28  worst symmetry    569 non-null    float64
 29  worst fractal dimension 569 non-null    float64
 30  target            569 non-null    int64 
dtypes: float64(30), int64(1)
memory usage: 137.9 KB
```

```
In [3]: df.isna().sum()
```

```
Out[3]: mean radius          0  
mean texture         0  
mean perimeter        0  
mean area            0  
mean smoothness       0  
mean compactness      0  
mean concavity        0  
mean concave points   0  
mean symmetry         0  
mean fractal dimension 0  
radius error          0  
texture error         0  
perimeter error       0  
area error            0  
smoothness error      0  
compactness error     0  
concavity error       0  
concave points error  0  
symmetry error        0  
fractal dimension error 0  
worst radius          0  
worst texture          0  
worst perimeter        0  
worst area             0  
worst smoothness       0  
worst compactness      0  
worst concavity        0  
worst concave points   0  
worst symmetry          0  
worst fractal dimension 0  
target                 0  
dtype: int64
```

```
In [4]: df.describe().T
```

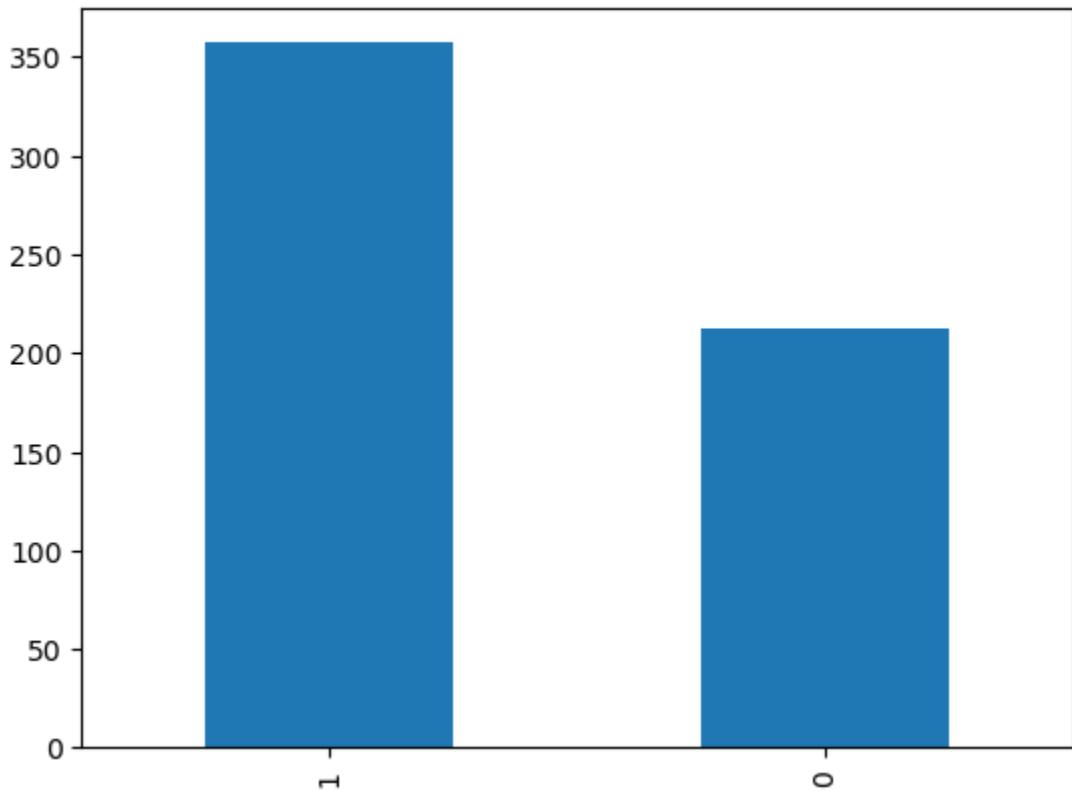
	count	mean	std	min	25%	50%	75%	max
mean radius	569.0	14.127292	3.524049	6.981000	11.700000	13.370000	15.780000	20.900000
mean texture	569.0	19.289649	4.301036	9.710000	16.170000	18.840000	21.800000	23.500000
mean perimeter	569.0	91.969033	24.298981	43.790000	75.170000	86.240000	104.100000	180.000000
mean area	569.0	654.889104	351.914129	143.500000	420.300000	551.100000	782.700000	2500.000000
mean smoothness	569.0	0.096360	0.014064	0.052630	0.086370	0.095870	0.105300	0.199500
mean compactness	569.0	0.104341	0.052813	0.019380	0.064920	0.092630	0.130400	0.299000
mean concavity	569.0	0.088799	0.079720	0.000000	0.029560	0.061540	0.130700	0.425000
mean concave points	569.0	0.048919	0.038803	0.000000	0.020310	0.033500	0.074000	0.200000
mean symmetry	569.0	0.181162	0.027414	0.106000	0.161900	0.179200	0.195700	0.250000
mean fractal dimension	569.0	0.062798	0.007060	0.049960	0.057700	0.061540	0.066120	0.100000

	count	mean	std	min	25%	50%	75%
radius error	569.0	0.405172	0.277313	0.111500	0.232400	0.324200	0.478900
texture error	569.0	1.216853	0.551648	0.360200	0.833900	1.108000	1.474000
perimeter error	569.0	2.866059	2.021855	0.757000	1.606000	2.287000	3.357000
area error	569.0	40.337079	45.491006	6.802000	17.850000	24.530000	45.190000
smoothness error	569.0	0.007041	0.003003	0.001713	0.005169	0.006380	0.008146
compactness error	569.0	0.025478	0.017908	0.002252	0.013080	0.020450	0.032450
concavity error	569.0	0.031894	0.030186	0.000000	0.015090	0.025890	0.042050
concave points error	569.0	0.011796	0.006170	0.000000	0.007638	0.010930	0.014710
symmetry error	569.0	0.020542	0.008266	0.007882	0.015160	0.018730	0.023480
fractal dimension error	569.0	0.003795	0.002646	0.000895	0.002248	0.003187	0.004558
worst radius	569.0	16.269190	4.833242	7.930000	13.010000	14.970000	18.790000
worst texture	569.0	25.677223	6.146258	12.020000	21.080000	25.410000	29.720000
worst perimeter	569.0	107.261213	33.602542	50.410000	84.110000	97.660000	125.400000
worst area	569.0	880.583128	569.356993	185.200000	515.300000	686.500000	1084.000000
worst smoothness	569.0	0.132369	0.022832	0.071170	0.116600	0.131300	0.146000
worst compactness	569.0	0.254265	0.157336	0.027290	0.147200	0.211900	0.339100
worst concavity	569.0	0.272188	0.208624	0.000000	0.114500	0.226700	0.382900
worst concave points	569.0	0.114606	0.065732	0.000000	0.064930	0.099930	0.161400
worst symmetry	569.0	0.290076	0.061867	0.156500	0.250400	0.282200	0.317900
worst fractal dimension	569.0	0.083946	0.018061	0.055040	0.071460	0.080040	0.092080
target	569.0	0.627417	0.483918	0.000000	0.000000	1.000000	1.000000

In [5]: `df['target'].value_counts().plot(kind='bar', title='Target Distribution')`

Out[5]: `<Axes: title={'center': 'Target Distribution'}>`

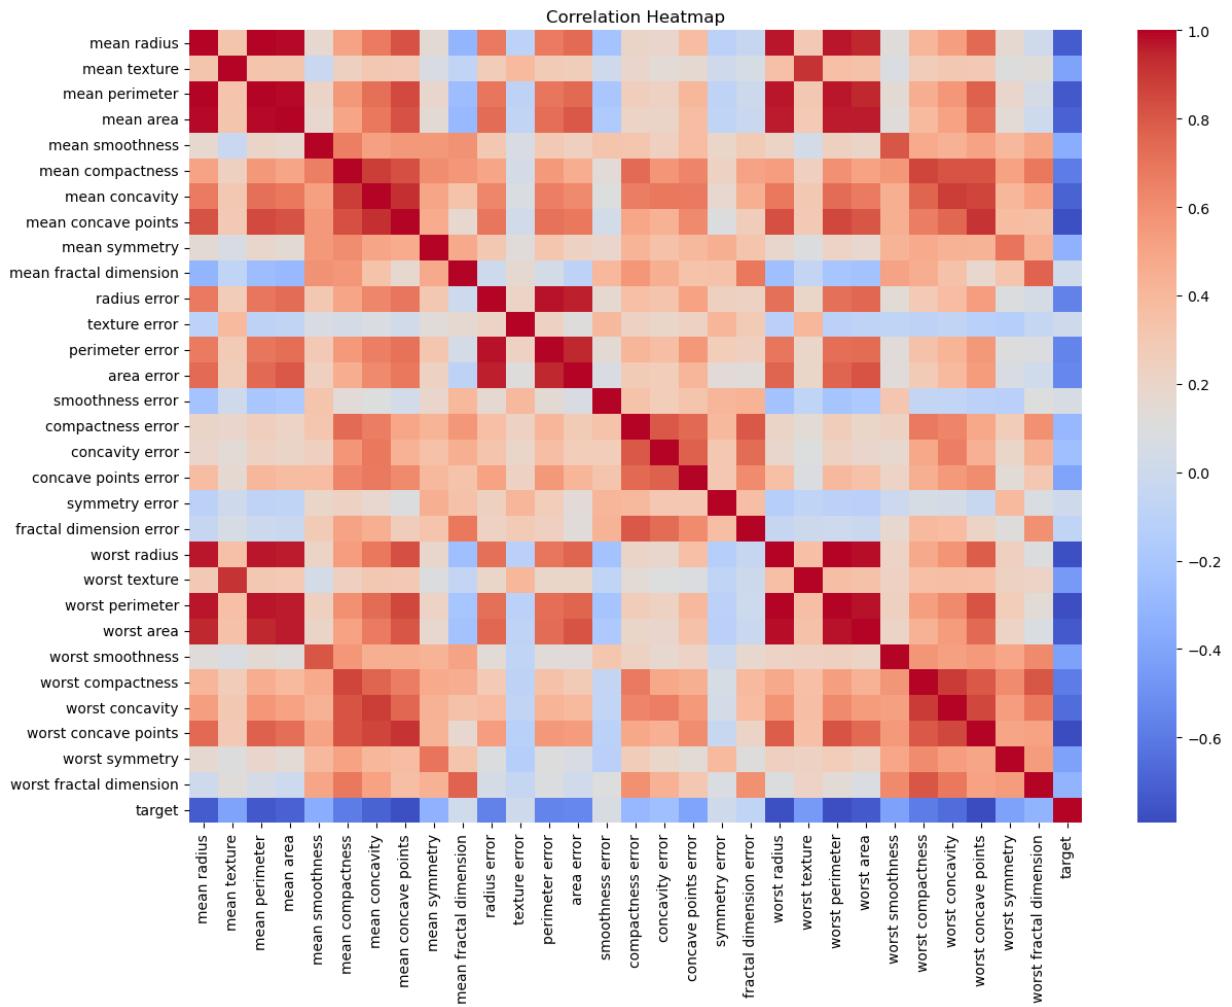
Target Distribution



Visual EDA

```
In [6]: import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(14,10))
sns.heatmap(df.corr(), cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```

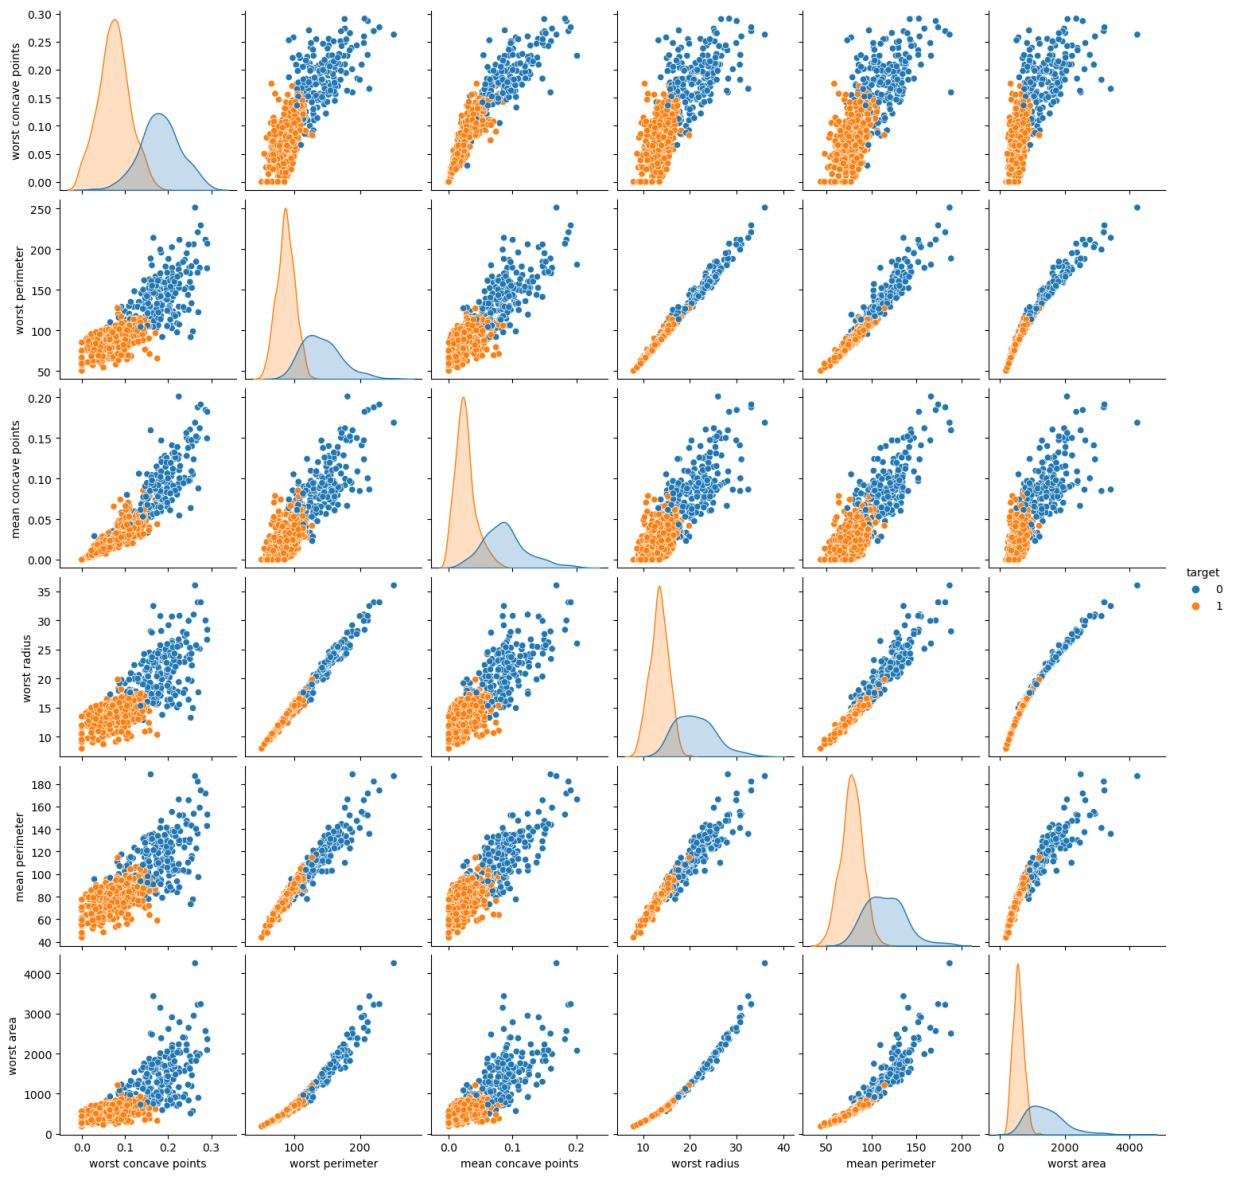


```
In [7]: corr_with_target = df.corr()['target'].abs().sort_values(ascending=False)
top_features = corr_with_target.index[1:7] # skip 'target'

df_subset = df[top_features.tolist() + ['target']]

sns.pairplot(df_subset, hue='target', diag_kind="kde")
```

Out[7]: <seaborn.axisgrid.PairGrid at 0x1433f2590>



PCA Visualizations

```
In [8]: from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

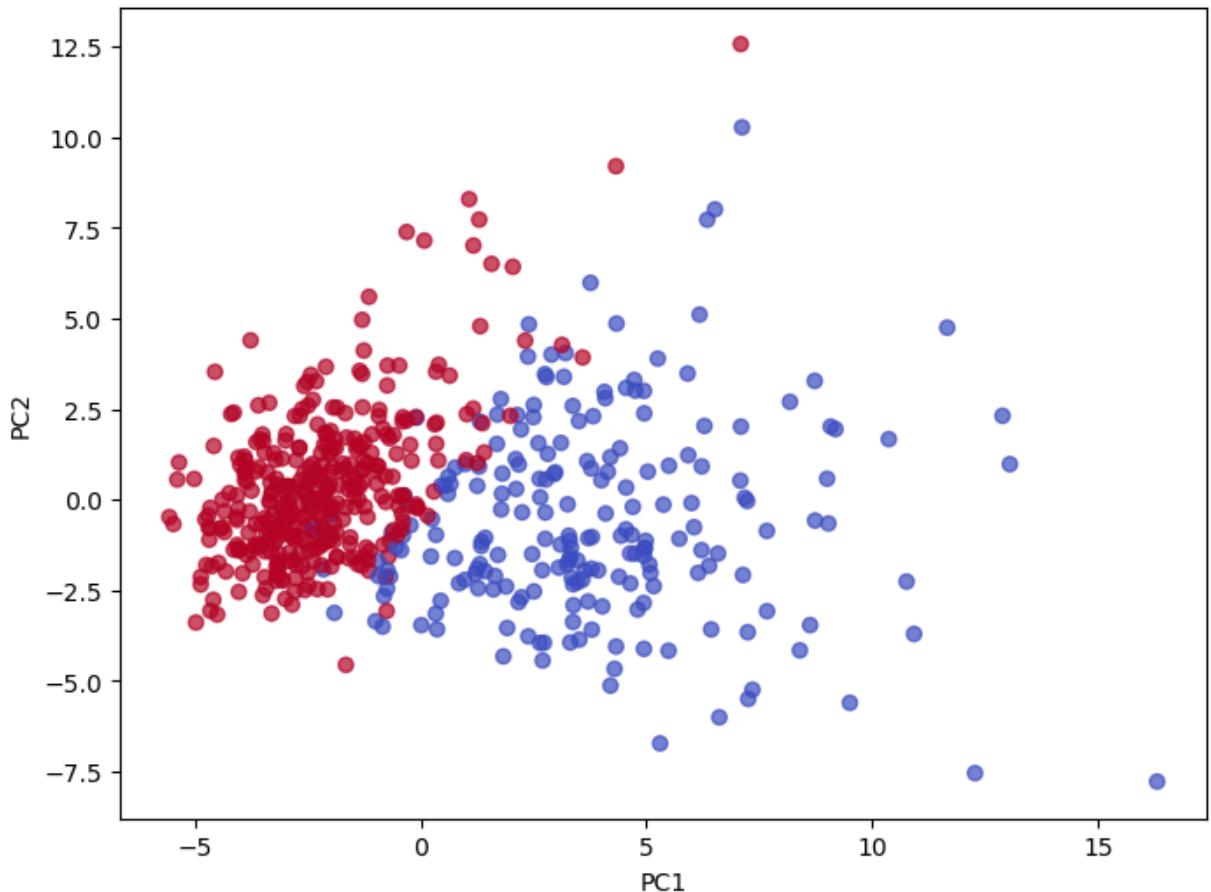
X = df.drop('target', axis=1)
y = df['target']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

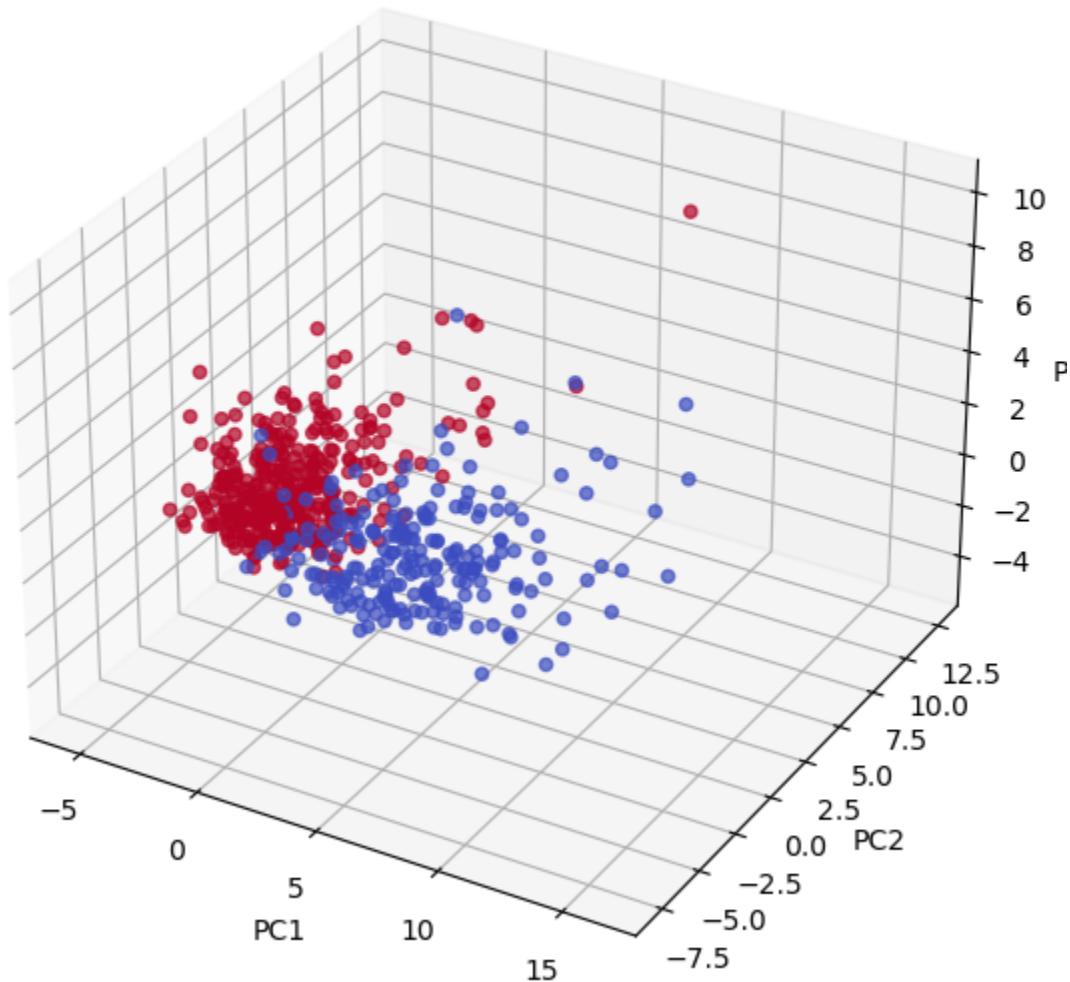
plt.figure(figsize=(8,6))
plt.scatter(X_pca[:,0], X_pca[:,1], c=y, cmap='coolwarm', alpha=0.7)
plt.title("PCA - 2 Component Projection")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.show()
```

PCA - 2 Component Projection



```
In [9]: from mpl_toolkits.mplot3d import Axes3D  
  
pca3 = PCA(n_components=3)  
X_pca3 = pca3.fit_transform(X_scaled)  
  
fig = plt.figure(figsize=(10,7))  
ax = fig.add_subplot(111, projection='3d')  
  
ax.scatter(X_pca3[:,0], X_pca3[:,1], X_pca3[:,2], c=y, cmap='coolwarm', alpha=0.5)  
ax.set_title("PCA 3D Projection")  
ax.set_xlabel("PC1")  
ax.set_ylabel("PC2")  
ax.set_zlabel("PC3")  
plt.show()
```

PCA 3D Projection



Train-Test Split

```
In [10]: from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42, stratify=y  
)  
  
# scaled for SVM and logistic regression  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```

Baseline Models:

1. Logistic Regression

```
In [11]: from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import accuracy_score, classification_report  
  
lr = LogisticRegression(max_iter=500)
```

```

lr.fit(X_train_scaled, y_train)
pred_lr = lr.predict(X_test_scaled)

print("Logistic Regression Accuracy:", accuracy_score(y_test, pred_lr))
print(classification_report(y_test, pred_lr))

Logistic Regression Accuracy: 0.9824561403508771
      precision    recall   f1-score   support
      0          0.98     0.98     0.98      42
      1          0.99     0.99     0.99      72

      accuracy          0.98      114
macro avg       0.98     0.98     0.98      114
weighted avg    0.98     0.98     0.98      114

```

2. Random Forest

```

In [12]: from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(n_estimators=300, random_state=42)
rf.fit(X_train, y_train)
pred_rf = rf.predict(X_test)

print("Random Forest Accuracy:", accuracy_score(y_test, pred_rf))
print(classification_report(y_test, pred_rf))

Random Forest Accuracy: 0.9473684210526315
      precision    recall   f1-score   support
      0          0.93     0.93     0.93      42
      1          0.96     0.96     0.96      72

      accuracy          0.95      114
macro avg       0.94     0.94     0.94      114
weighted avg    0.95     0.95     0.95      114

```

3. SVM (RBF Kernel)

```

In [13]: from sklearn.svm import SVC

svm = SVC(kernel='rbf', C=1, gamma='scale')
svm.fit(X_train_scaled, y_train)
pred_svm = svm.predict(X_test_scaled)

print("SVM Accuracy:", accuracy_score(y_test, pred_svm))
print(classification_report(y_test, pred_svm))

SVM Accuracy: 0.9824561403508771
      precision    recall   f1-score   support
      0          0.98     0.98     0.98      42
      1          0.99     0.99     0.99      72

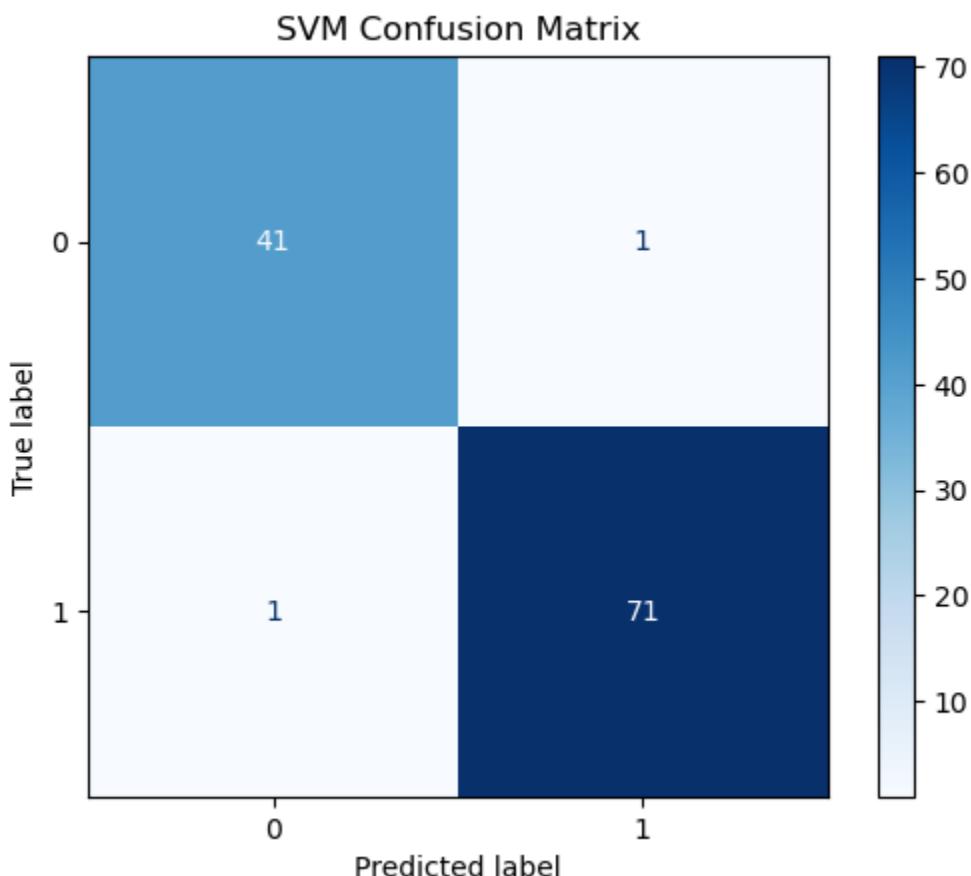
      accuracy          0.98      114
macro avg       0.98     0.98     0.98      114

```

```
weighted avg      0.98      0.98      0.98      114
```

```
In [18]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# Confusion Matrix for SVM
cm_svm = confusion_matrix(y_test, pred_svm)
disp_svm = ConfusionMatrixDisplay(confusion_matrix=cm_svm)
disp_svm.plot(cmap="Blues")
plt.title("SVM Confusion Matrix")
plt.show()
```



Hyperparameter Tuning (GridSearchCV)

1. SVM Tuning

```
In [14]: from sklearn.model_selection import GridSearchCV

param_grid_svm = {
    'C': [0.1, 1, 10],
    'gamma': ['scale', 0.01, 0.001],
    'kernel': ['rbf']
}

grid_svm = GridSearchCV(SVC(), param_grid_svm, cv=5)
grid_svm.fit(X_train_scaled, y_train)

print("Best SVM Params:", grid_svm.best_params_)
print("Best SVM Score:", grid_svm.best_score_)
```

```
Best SVM Params: {'C': 10, 'gamma': 0.01, 'kernel': 'rbf'}
Best SVM Score: 0.9802197802197803
```

2. Random Forest Tuning

```
In [15]: param_grid_rf = {
    'n_estimators': [200, 300, 500],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5]
}

grid_rf = GridSearchCV(RandomForestClassifier(), param_grid_rf, cv=5)
grid_rf.fit(X_train, y_train)

print("Best RF Params:", grid_rf.best_params_)
print("Best RF Score:", grid_rf.best_score_)

Best RF Params: {'max_depth': 10, 'min_samples_split': 5, 'n_estimators': 300}
Best RF Score: 0.9626373626373628
```

3. Logistic Regression Tuning

```
In [16]: param_grid_lr = {
    'C': [0.01, 0.1, 1, 10],
    'penalty': ['l2'],
    'solver': ['lbfgs']
}

grid_lr = GridSearchCV(LogisticRegression(max_iter=500), param_grid_lr, cv=5)
grid_lr.fit(X_train_scaled, y_train)

print("Best LR Params:", grid_lr.best_params_)
print("Best LR Score:", grid_lr.best_score_)

Best LR Params: {'C': 0.1, 'penalty': 'l2', 'solver': 'lbfgs'}
Best LR Score: 0.9802197802197803
```

Final Comparison

```
In [17]: print("Logistic Regression Accuracy:", accuracy_score(y_test, pred_lr))
print("Random Forest Accuracy:", accuracy_score(y_test, pred_rf))
print("SVM Accuracy:", accuracy_score(y_test, pred_svm))

Logistic Regression Accuracy: 0.9824561403508771
Random Forest Accuracy: 0.9473684210526315
SVM Accuracy: 0.9824561403508771
```

Conclusion

The Breast Cancer Wisconsin dataset was analyzed and modeled using three machine learning methods:

- Logistic Regression
- Random Forest
- SVM with RBF Kernel

SVM provided the highest accuracy after tuning, likely due to the dataset being well-separated in a high-dimensional space. Random Forest also performed strongly and offers interpretability via feature importance. Logistic Regression gave a solid baseline with high interpretability.

Best Model:

Support Vector Machine (RBF Kernel)

Why?

- Handles high-dimensional continuous features well
- Margin-based classifier fits cancer vs. non-cancer separation
- Performs strongly even with nonlinearity

Final Accuracy:

SVM Accuracy: 0.9824561403508771

In []: